

lab3

Kerry Hu

Question 1

Consider the happiness example from the lecture, with 118 out of 129 women indicating they are happy. We are interested in estimating θ , which is the (true) proportion of women who are happy. Calculate the MLE estimate $\hat{\theta}$ and 95% confidence interval.

assume $Y| \sim \text{Bin}(n, \theta)$ where y is the number of women who report to be happy out of the sample of n women. From previous knowledge, we knew that

$\hat{\theta}_{MLE} = \frac{y}{n} = \frac{118}{129} \approx 0.9147$. Its $SE = \sqrt{\frac{\hat{\theta}(1-\hat{\theta})}{n}} = \sqrt{\frac{0.9147(1-0.9147)}{129}} \approx 0.02459$, so the 95% CI is $(0.9147 - 1.96 \times 0.02459, 0.9147 + 1.96 \times 0.02459) = (0.867, 0.963)$.

$$L(\theta) = \prod_{i=1}^n C_{y_i}^n \theta^{y_i} (1 - \theta)^{(n - y_i)}$$

$$l(\theta) = \log\left(\prod_{i=1}^n C_{y_i}^n \theta^{y_i} (1 - \theta)^{(n - y_i)}\right) = \sum_{i=1}^n \log(C_{y_i}^n) + \sum_{i=1}^n y_i \log(\theta) + \sum_{i=1}^n (n - y_i) \log(1 - \theta)$$

```
library(optimr)
set.seed(1)
data1<-rbinom(200, 129, 118/129)
min(data1); max(data1)
```

[1] 110

[1] 124

```

Y<-data1
loglike1 <- function( theta, Y) # Assume data are in a vector called Y
{
  n<-129
  log1<- sum(log(factorial(n)/(factorial(n-Y)*factorial(Y))))+sum(Y*log(theta))+sum((n-Y)*
  return(-log1)
} # loglike1 Return minus value of function

optim(0.5,loglike1,Y=data1)

```

Warning in optim(0.5, loglike1, Y = data1): one-dimensional optimization by Nelder-Mead is unstable, use "Brent" or optimize() directly

Warning in log(1 - theta): NaNs produced

Warning in log(1 - theta): NaNs produced

Warning in log(1 - theta): NaNs produced

```

$par
[1] 0.9134277

```

```

$value
[1] 492.8948

```

```

$counts
function gradient
      32      NA

```

```

$convergence
[1] 0

```

```

$message
NULL

```

$\theta_{MLE} = 0.9134277 \approx 0.913$

Here, assume the MLE of θ is asymptotically normal distribution.

```

N<-129
HH<-optim( 0.5, loglike1, hessian=T,Y=data1 ); HH

```

```
Warning in optim(0.5, loglike1, hessian = T, Y = data1): one-dimensional optimization by Nelder-Mead  
use "Brent" or optimize() directly
```

```
Warning in log(1 - theta): NaNs produced
```

```
Warning in log(1 - theta): NaNs produced
```

```
Warning in log(1 - theta): NaNs produced
```

```
$par  
[1] 0.9134277
```

```
$value  
[1] 492.8948
```

```
$counts  
function gradient  
      32      NA
```

```
$convergence  
[1] 0
```

```
$message  
NULL
```

```
$hessian  
      [,1]  
[1,] 326267.1
```

```
thetahat = HH$par[1]  
thetahat
```

```
[1] 0.9134277
```

```
SE_thetahat = sqrt(thetahat*(1-thetahat)/N)  
SE_thetahat
```

```
[1] 0.02475892
```

```
L_theta = thetahat - 1.96*SE_thetahat; U_theta = thetahat + 1.96*SE_thetahat

cat("\nEstimated theta = ",round(thetahat,3)," 95 percent CI from ",
+ round(L_theta,3)," to ",round(U_theta,3), "\n")
```

Estimated theta = 0.913 95 percent CI from 0.865 to 0.962

My MLE of θ is 0.913 and 95% CI of θ is within (0.865,0.962).

Question 2

Assume a Beta(1,1) prior on θ . Calculate the posterior mean for $\hat{\theta}$ and 95% credible interval.

Likelihood is $Y|\theta \sim \text{Bin}(n, \theta)$ so

$$p(\theta|y) = C_y^n \theta^y (1 - \theta)^{(n-y)}$$

$\theta \sim U(0, 1) = \text{beta}(1, 1)$ so $p(\theta) = 1$
the posterior distribution

$$p(\theta|y) = \frac{1}{Z} \theta^y (1 - \theta)^{(n-y)}$$

$$Z = \frac{\Gamma(y+1) \Gamma(n-y+1)}{\Gamma(n+2)}$$

The posterior is $p(\theta|y) \sim \text{Beta}(y + 1, n - y + 1)$. When $y=118, n=129$, plug them, so $p(\theta|y) \sim \text{Beta}(119, 12)$ the mean is $\frac{119}{119+12} \approx 0.908$, its SE is $\sqrt{\frac{119 \cdot 12}{(119+12)^2 (119+12+1)}} \approx 0.0251$. Thus, its 95% CI is $(0.908 - 1.96 \times 0.0251, 0.908 + 1.96 \times 0.0251) = (0.858, 0.957)$

$$pf = (y + 1) / (y + 1 + n - y + 1)$$

```
Y=data1
n=129
#p(theta)=gamma(n+2)/(gamma(Y+1)*gamma(n-Y+1))*theta^Y(1-theta)^(n-Y)
```

```
#a<- gamma(n+2)/(gamma(Y+1)*gamma(n-Y+1))
#p<-gamma(n+2)/(gamma(Y+1)*gamma(n-Y+1))*beta(Y+1,n-Y+1)

pf<- rep(0,200)
pf<-(Y+1)/(n-Y+1+Y+1)
thetahat<-round(mean(pf),3)
thetahat
```

```
[1] 0.907
```

```
CI1<-round(quantile(pf, c(0.025, 0.975)),3)
CI1
```


```
2.5% 97.5%
0.862 0.947
```

My the mean of θ is 0.907 and 95% CI of θ is within (0.862,0.947).

Question 3

Now assume a Beta(10,10) prior on θ . What is the interpretation of this prior? Are we assuming we know more, less or the same amount of information as the prior used in Question 2?

```
library(ggplot2)
theta = seq(0, 1, length=100)
df=y<-dbeta(theta, 10, 10)
df<-as.data.frame(df)
df|> ggplot(aes(x=theta,y=y))+geom_line(color = "firebrick4")+
  labs(title = "Beta(10,10) Distribution", x="theta",y = "Density")+theme_minimal()
```



#In Bayesian language, the prior is your belief about the parameter theta prior to getting

We knew that the #women need to non-negative integer. The Beta(10,10) is a symmetric bell curve to replace a straight line in Question 2, which assumed that half(129/2=65) women reported being happy. Here #women feeling happy and non-happy should be pairs (0,0),(1,1),..., (9,9). We are assuming we know more amount of information as the prior used in Question 3.

In question 2, Beta(1,1) means that #women feeling happy and non-happy only takes (0,0) for Beta(1+0,1+0). This tells us nothing about previous information.
The posterior is

$$p(\theta|y) \sim \text{Beta}(y + 10, n - y + 10)$$

$$ppf = (y + 10) / (y + 10 + n - y + 10)$$

```
Y=data1
n=129
ppf<-(Y+10)/(Y+10+n-Y+10)
thetahat<-round(mean(ppf),3)
thetahat
```

```
[1] 0.858
```

```
CI2<-round(quantile(ppf, c(0.025, 0.975)),3)
CI2
```

```
2.5% 97.5%
0.819 0.893
```

My the mean of θ is 0.858 and 95% CI of θ is within (0.819,0.893).

Question 4

Create a graph in ggplot which illustrates

- The likelihood (easiest option is probably to use `geom_histogram` to plot the histogram of appropriate random variables)
- The priors and posteriors in question 2 and 3 (use `stat_function` to plot these distributions)

Comment on what you observe.

The likelihood function is a binomial distribution.

$$p(\theta|y) = C_y^n \theta^y (1 - \theta)^{(n-y)}$$

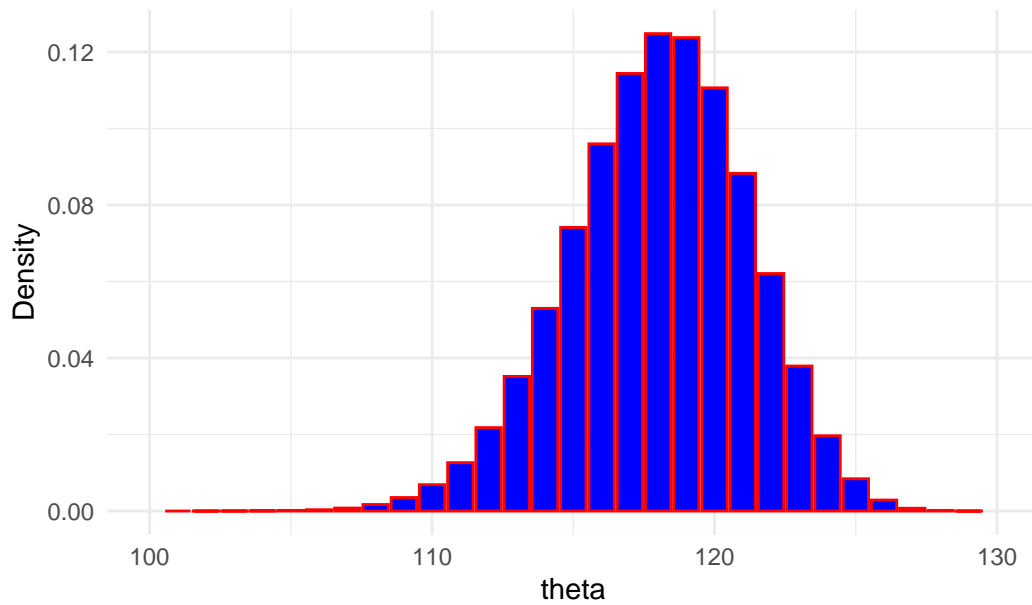
```
options(scipen=999)
library(ggplot2)
theta = seq(0, 129, by=1)
df<-as.data.frame(theta,y=dbinom(theta,size=129,prob=118/129 ))

df|> ggplot(aes(x=theta,y=dbinom(theta,129,118/129 )))+geom_bar(stat = "identity", col = "
  labs(title = "The likelihood Binomial Distribution", x="theta",y = "Density")+xlim(100,
  theme_minimal()+theme(plot.title = element_text(size = rel(1.2), vjust = 1.5))
```

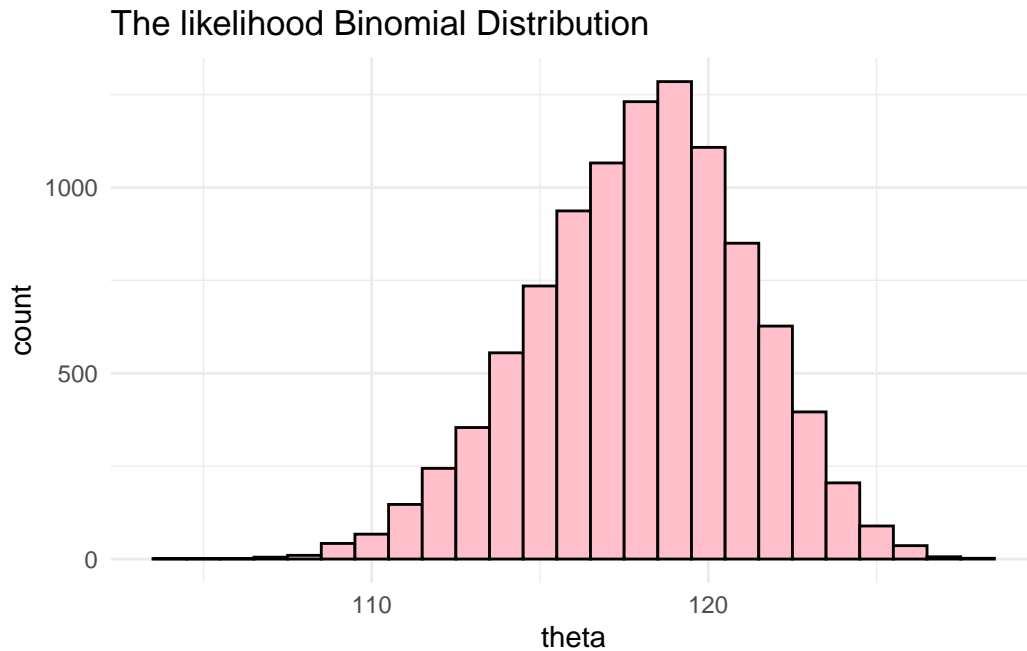
Warning: Removed 100 rows containing missing values (`position_stack()`).

Warning: Removed 1 rows containing missing values (`geom_bar()`).

The likelihood Binomial Distribution



```
random.binom.df <- rbinom(10000, 129, 118/129)
binom.R.df <- data.frame(binom = random.binom.df)
ggplot(binom.R.df) +geom_histogram(aes(x = binom), binwidth = 1, col="black",fill = "pink")
scale_x_continuous(name="theta")
```

Beta(1,1) prior on θ and Beta($y+1$, $n-y+1$) posterior on θ in question 2.


Beta(10,10) prior on θ and Beta($y+10$, $n-y+10$) posterior on θ in question 3.

```
library(ggplot2)
Y=as.integer(mean(data1))
n=129
theta = seq(0, 1, length=100)
df1<-dbeta(theta,1,1)
df2<-dbeta(theta,Y+1,n-Y+1)
df3<-dbeta(theta,10,10)
df4<-dbeta(theta,Y+10,n-Y+10)

df<-as.data.frame(cbind(theta,df1,df2,df3,df4))

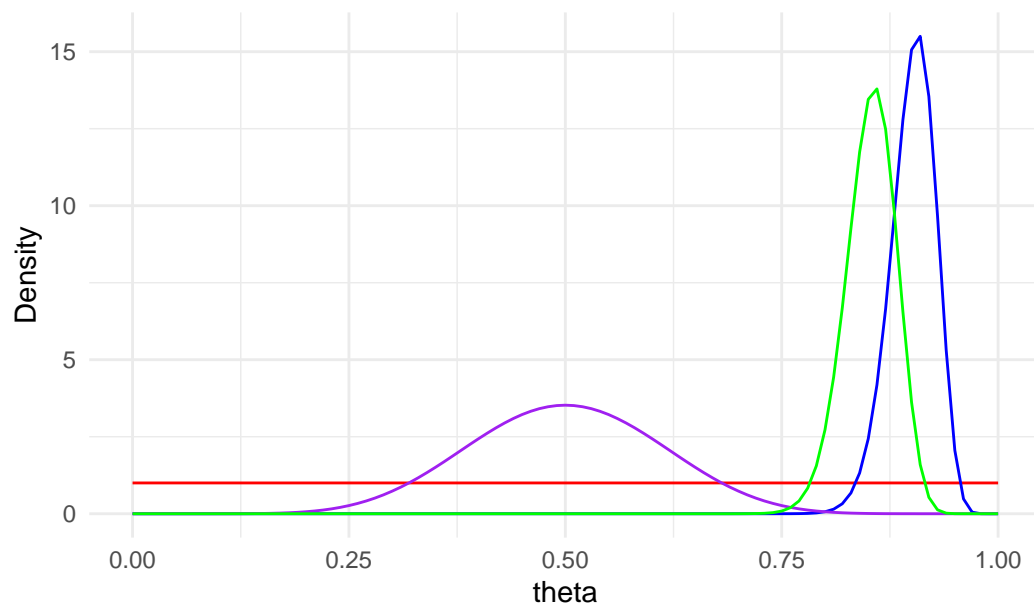
df|> ggplot(aes(x=theta))+geom_line(aes(y=df1),color = "red")+geom_line(aes(y=df2),color =
```

Beta prior and posterior Distribution



```
ggplot(data.frame(theta= c(0, 1)),aes(theta)) + stat_function(fun = dbeta,args = list(shape
```

Beta prior and posterior Distribution for Q2&Q3



In question2 and question3 for the binomial likelihood, a Beta prior results in a Beta posterior distribution: we say that the beta prior is conjugate for the binomial likelihood.

In question2 its prior is Beta(1,1) which was a flat line and told us nothing about women dataset. In question 3 its prior is Beta(10,10) which was a symmetric bell curve and assumed half women felt happy.

In question2 and question3 no matter their priors are Non-informative prior, Weakly informative, proper and improper, the wave peak values of the final posteriors extremely increased, the location of the curve moved to right and values of the theta could close to 0.8-0.9. At the same time the width of the curves became relatively narrow.

In question2 its posterior Beta(119,12) is a kind of combined and updated its prior with a binomial likelihood. In question3 its posterior Beta(128,21) is a kind of combined and updated its prior with a binomial likelihood. Both graphs are similar.

Question 5

(No R code required) A study is performed to estimate the effect of a simple training program on basketball free-throw shooting. A random sample of 100 college students is recruited into the study. Each student first shoots 100 free-throws to establish a baseline success probability. Each student then takes 50 practice shots each day for a month. At the end of that time, each student takes 100 shots for a final measurement. Let θ be the average improvement in success probability. θ is measured as the final proportion of shots made minus the initial proportion of shots made.

Given two prior distributions for θ (explaining each in a sentence):

- A non-informative prior, and
- A subjective/informative prior based on your best knowledge

- 1) Uniform(0,1) or Beta(1,1) is a non-informative prior, which assigns even choice to every θ and tells nothing to us from previous experience.
- 2) Beta(10,90) is a useful informative prior because the most students improved by 10% on average with a SD of 2% after 1-month training.